

Pintos Virtual Memory

Daniel Chiu

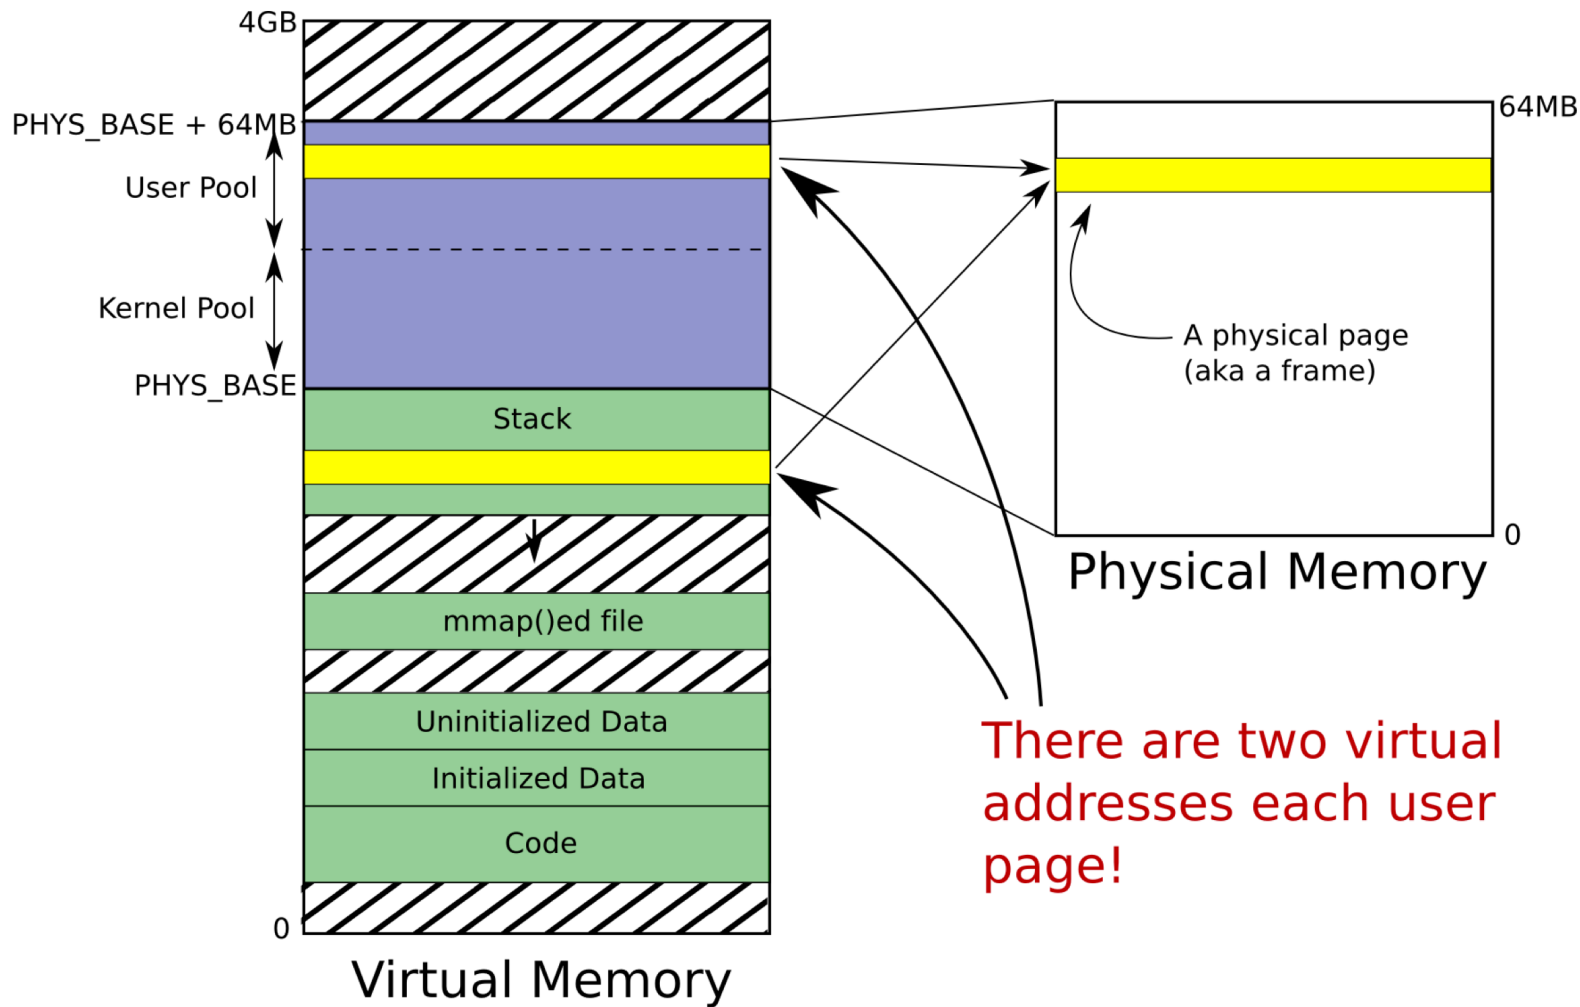
Slides adapted from previous quarters

Overview: Project Components

- High Level Goal: Implement Virtual Memory
- Page Table Management
 - Page Fault Handler
 - Interacts with most of the rest of the system
 - Eviction Policy
 - Approximation of LRU (e.g. “clock algorithm”)
- Swap Disk Management
 - Page{in,out} evicted pages
- Stack Growth
- Memory Mapped Files (including lazy loading executables)

First: Memory Layout

- All user virtual addresses have corresponding kernel virtual addresses
- The hardware page table (discussed in lecture) allows the MMU to map addresses automatically
- You will need to keep around additional information
 - e.g. keep track of where to fetch pages that aren't in memory
 - Explained in a few more slides



Pintos Page Tables

- Base implementation already creates basic page directory & page table structure mappings. (Look at `paging_init()` in `init.c`)
- e.g. user wants a new upage at `vaddr`:
 - `palloc_get_page(PAL_USER)` returns a page
 - Register with `pagedir_set_page()`
- Layout is very similar to the one discussed in lecture, except:
 - Only two levels of page tables
 - 4 byte page table entries
- Note that hardware bits (present, writable, user, accessed, dirty) are **PER PAGE TABLE ENTRY**, not per physical frame!

Handling Virtual Memory

- Page Fault Handler

- Verify that the access is legal (read/write)
- If page has been evicted, load from disk/swap
- Else, trigger stack growth if necessary
- Paging something in might require paging something else out
- Beware of race conditions while accessing data structures
- You may need to “pin” pages to prevent them from being evicted
 - Don't pin for too long, but
 - Don't worry about running out of pages to pin (you can panic the kernel)

- Global Data Structure: **Frame Table**

- Keeps track of the *physical frames* that are allocated/free

- Per-Process Data Structure: **Supplemental Page Table**

- Keeps track of supplemental data about each page
- i.e. location of data (frame/disk/swap); pointer to corresponding kernel virtual address; list of aliases, etc.

Stack Growth

- As part of the page fault handler, devise a heuristic to identify valid stack accesses that require stack growth
- Valid stack accesses may cause the following page faults:
 - PUSH: 4 bytes below %esp
 - PUSHA: 32 bytes below %esp
 - Other accesses: anywhere between %esp and PHYS_BASE
 - Accesses to stack pages that have been paged out
 - Accesses above the stack pointer after %esp has been decremented (i.e. with SUB \$n %esp)
- All pages on the stack (except the first one) must be allocated lazily
- You should limit stack size to a constant as most OSes do
 - Why is this a good idea?

Memory Mapped Files

- Implement the `mmap()` and `munmap()` system calls
- Per-Process Data Structure: **mmap Page Table**
 - Keeps track of file -> page mappings for the process
- `mmap()`ed pages must be loaded from disk lazily
- `mmap()` should fail (return -1) if:
 - The size of the file is zero bytes
 - The range of pages mapped overlaps with existing mapped pages
- All mappings are implicitly unmapped when a process exits

Page Replacement Mechanism

- Eviction Policy

- Approximation of LRU replacement (e.g. “clock” algorithm)
- Pages that can be deallocated without writing to swap/disk
 - Read only (code) pages
 - Clean data pages
 - Clean mmap()ed pages
- Pages that must be written out to swap
 - Stack pages
 - Dirty data pages
- Write dirty mmap()ed pages to their corresponding file on disk

- Global Data Structure: **Swap Table**

- Keeps track of the *swap slots* that are allocated/free

Implementation Order


Frame Table: allocate frames (physical pages)



Supplemental Page Table: handle page faults



Stack growth, mmap()ed files, lazy executable loading [work in parallel]



Eviction/paging

But these are all interdependent! It WON'T work to design them separately and try to piece everything together later!

Useful Code

- Page Allocator (threads/palloc.{h,c})
 - Allows you to obtain a virtual address for a physical frame
 - Pages can be allocated from either the user pool or the kernel pool
- Page Directory (userprog/pagedir.{h,c}, threads/pte.h)
 - Interface to the x86 hardware page directory
 - Allows you to perform operations on the hardware page tables for each process
 - pte.h gives access to hardware bits (p/w/u/a/d)
- General data structures
 - Hash Table (lib/kernel/hash.{h,c})
 - Bitmap (lib/kernel/bitmap.{h,c})
 - List (lib/kernel/list.{h,c}) [you should already know this one!]
- Refer to Appendix A of the Pintos Reference Guide for a more elaborate description of the available APIs

No Extra Credit

- In previous years there was extra credit available for this assignment
- There is NO extra credit being offered for this assignment (or for assignment 4) despite what it says in the assignment handout!

Good Luck!

- Any Questions?